


Training Loss Curves — Binary DLGN vs Ternary PST-DTLGN

Vlarge (192K neurons, 4L)



Huge (512K neurons, 4L)

