


$R_F$ 's optimum IS the dead gate  
(accuracy rises — but by pruning, not interpretability)



The L1 drop is coefficient-scale shrink + gate death,  
NOT spectral-shape sparsification

